

A decorative graphic on the left side of the slide, consisting of a network of white lines and circles on a blue gradient background, resembling a circuit board or a neural network.

A Concurrent Implementation of MapReduce

CS5600 FINAL PROJECT

DECEMBER 16, 2021

BY: MITCHELL LAWSON

MOTIVATION

- Redemption for the dreaded File Systems assignment
- Redemption from the concurrent Hash Table assignment
- Learn more about concurrency
- Attempt something hard and very interesting
- Prepare for Scalable Distributed Systems → MapReduce, created by engineers at Google for large scale parallel data processing
 - It makes programming on large-scale clusters easier for developers by alleviating the worries of managing parallelism, crashes and other issues associated with clusters of machines.

PROJECT DETAILS

- Project is defined in the OSTEP repository:
 - <https://github.com/remzi-arpacidusseau/ostep-projects/tree/master/concurrency-mapreduce>.
- A simplification that is made for this project as suggested by the outline given is creating MapReduce for a single machine. There are 3 learning objectives for this personal end of term project.
 1. Learn about and understand the nature of the MapReduce paradigm.
 2. Implement a correct and efficient version of MapReduce that utilizes threads.
 3. Gain more experience and familiarity creating concurrent code.
- Will be used/implemented in the context of word counts

WELL... WHAT ACTUALLY IS MAPREDUCE?

- MapReduce can perform distributed and parallel computations using large datasets across a large number of nodes
- A MapReduce job usually splits the input datasets and then process each of them independently by the Map tasks in a completely parallel manner... BUT in my implementation this is done in a concurrent manner

THINK OF THE FOLLOWING COMPONENTS

1. Mapper

- Some logic is applied to "n" data blocks spread across different node.
- The data blocks in my implementation are words and the nodes are files.

2. Shuffle and Sort

- The output of the mapper then gets sorted
- Repeated values get removed and the values get grouped together based on keys
- This was done using quick sort and a concurrent hash table/ map

3. Reducer

- The output from the previous step is used in this phase.
- "Consolidates" the outputs and creates the finalized product.
- Counts the words in my case.

4. Combiner (Optional)

- Provides performance increase by combining inputs from the same node.

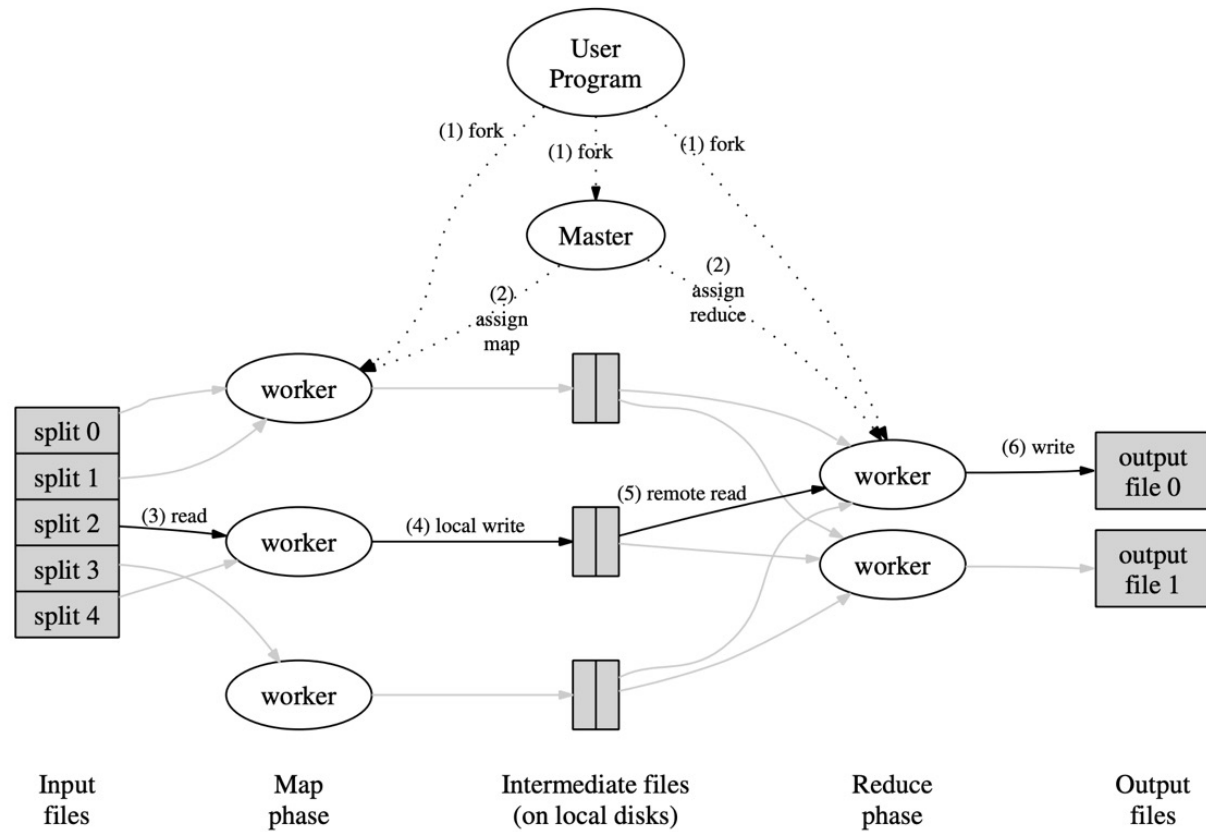


Figure 1: Execution overview

THE PROVIDED INTERFACE

```
Concurrent_MapReduce - mapreduce.h
1  #ifndef __mapreduce_h__
2  #define __mapreduce_h__
3
4  // Different function pointer types used by MR
5  typedef char *(*Getter)(char *key, int partition_number);
6  typedef void (*Mapper)(char *file_name);
7  typedef void (*Reducer)(char *key, Getter get_func, int partition_number);
8  typedef unsigned long (*Partitioner)(char *key, int num_partitions);
9
10 // External functions: these are what you must define
11 void MR_Emit(char *key, char *value);
12
13 unsigned long MR_DefaultHashPartition(char *key, int num_partitions);
14
15 void MR_Run(int argc, char *argv[],
16             Mapper map, int num_mappers,
17             Reducer reduce, int num_reducers,
18             Partitioner partition);
19
20 #endif // __mapreduce_h__
21
```


WORD COUNTING

```
12 void Map(char *file_name) {
13     FILE *fp = fopen(file_name, "r");
14     assert(fp != NULL);
15
16     char* line = NULL;
17     size_t size = 0;
18     while (getline(&line, &size, fp) != -1) {
19         char *token, *dummy = line;
20         while ((token = strsep(&dummy, "\t\n\r")) != NULL) {
21             MR_Emit(token, "1");
22         }
23     }
24
25     fclose(fp);
26 }
27
28 void Reduce(char *key, Getter get_next, int partition_number) {
29     int count = 0;
30     char *value;
31     value = get_next(key, partition_number);
32     while (value != NULL) {
33         count++;
34         value = get_next(key, partition_number);
35     }
36     printf("%s %d\n", key, count);
37 }
38
39 int main(int argc, char *argv[]) {
40     MR_Run(argc, argv, Map, num_mappers: 10, Reduce, num_reducers: 10, MR_DefaultHashPartition);
41 }
42
```


TESTS

- There are many different areas of my implementation that I wanted to test, however due to time constraints I had to narrow it down to two sets of testing involving...
- Ensuring the correct output
 1. 1 file containing non-repeating words.
 2. 1 file containing repeating and non-repeating words.
 3. Multiple files containing repeating, non-repeating words and overlapping words.
- The timing of different file contents with my Reducer
 1. Best Case: 1 file all the same word
 2. Average Case: 1 file with a real-world example from free published text.
 3. Worst Case: All different words.

RESULTS & ANALYSIS

Timing tests:

1. Best Case: 1 file all the same word.
2. Average Case: 1 file with a real-world example from free published text.
3. Worst Case: All different words.

Test	Time (s)
1 (Best)	0.025456, 0.026044, 0.024939
2 (Average)	0.223773, 0.199620, 0.192057
3 (Worst)	1.093060, 0.953125, 0.912240

- The timing of different file contents with my Reducer

1. Best Case: 3 files, 10 mapping threads
2. Average Case: 3 files, 3 mapping threads
3. Worst Case: 3 files, 2 mapping threads

Test	Time (s)
1 (Best)	1.005772
2 (Average)	1.018589
3 (Worst)	1.031590

- The timing of different file contents with my Reducer

1. Best Case: 3 files, 2 reducing threads
2. Average Case: 3 files, 3 reducing threads
3. Worst Case: 3 files, 10 reducing threads

Test	Time (s)
1 (Best)	0.530489
2 (Average)	0.878015
3 (Worst)	1.047338

BENEFITS

- Fault Tolerance
- Resilience
- Fast
- Parallel Processing
- Availability
- Scalability
- Cost Effective